



Experiment No.5 Diffusion Models

Name : Gourav Sable

PRN : 202302040019

Implement a basic diffusion model for image generation by simulating the process of adding and removing noise. Use a pre-trained diffusion model to generate images from text prompts and compare results, highlighting model strengths, limitations, and ethical considerations.

Topic : Movie poster Generator

Link : https://colab.research.google.com/drive/1gGXoj6R_2luVNOkmu0WvmyJgaRk3Vbqb?usp=sharing

```
%pip install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu118
%pip install diffusers transformers accelerate
```

```
Looking in indexes: https://download.pytorch.org/whl/cu118
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: torchvision in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: torchaudio in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-nccl-cu12==2.27.3 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages
```

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages/MarkupSafe-2.0-py3.12.egg

Requirement already satisfied: diffusers in /usr/local/lib/python3.12/dist-packages/diffusers-0.20.0-py3.12.egg

Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages/transformers-4.30.2-py3.12.egg

Requirement already satisfied: accelerate in /usr/local/lib/python3.12/dist-packages/accelerate-0.17.0-py3.12.egg

Requirement already satisfied: importlib_metadata in /usr/local/lib/python3.12/dist-packages/importlib_metadata-4.10.0-py3.12.egg

Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages/filelock-3.4.3-py3.12.egg

Requirement already satisfied: huggingface-hub>=0.34.0 in /usr/local/lib/python3.12/dist-packages/huggingface_hub-0.34.0-py3.12.egg

Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages/numpy-1.23.5-py3.12.egg

Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.12/dist-packages/regex-2023.10.1-py3.12.egg

Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages/requests-2.31.0-py3.12.egg

Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.12/dist-packages/safetensors-0.3.1-py3.12.egg

Requirement already satisfied: Pillow in /usr/local/lib/python3.12/dist-packages/Pillow-9.4.0-py3.12.egg

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages/packaging-20.9-py3.12.egg

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages/pyyaml-5.4-py3.12.egg

Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages/tokenizers-0.22.0-py3.12.egg

Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.12/dist-packages/tqdm-4.60.0-py3.12.egg

Requirement already satisfied: psutil in /usr/local/lib/python3.12/dist-packages/psutil-3.4.1-py3.12.egg

Requirement already satisfied: torch>=2.0.0 in /usr/local/lib/python3.12/dist-packages/torch-2.0.0-py3.12.egg

Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.12/dist-packages/fsspec-2023.5.0-py3.12.egg

Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.12/dist-packages/typing_extensions-3.7.4.3-py3.12.egg

Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.12/dist-packages/hf_xet-1.1.3-py3.12.egg

Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages/setuptools-65.5.0-py3.12.egg

Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages/sympy-1.13.3-py3.12.egg

Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-packages/networkx-2.9.1-py3.12.egg

Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages/jinja2-3.1.2-py3.12.egg

Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages/nvidia_cuda_nvrtc_cu12-12.6.77-py3.12.egg

Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages/nvidia_cuda_runtime_cu12-12.6.77-py3.12.egg

Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages/nvidia_cuda_cupti_cu12-12.6.80-py3.12.egg

Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages/nvidia_cudnn_cu12-9.10.2.21-py3.12.egg

Requirement already satisfied: nvidia-cUBLAS-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages/nvidia_cUBLAS_cu12-12.6.4.1-py3.12.egg

Requirement already satisfied: nvidia-cUFFT-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages/nvidia_cUFFT_cu12-11.3.0.4-py3.12.egg

Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages/nvidia_curand_cu12-10.3.7.77-py3.12.egg

Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages/nvidia_cusolver_cu12-11.7.1.2-py3.12.egg

Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages/nvidia_cusparse_cu12-12.5.4.2-py3.12.egg

Requirement already satisfied: nvidia-cusparseLTLT-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages/nvidia_cusparseLTLT_cu12-0.7.1-py3.12.egg

Requirement already satisfied: nvidia-NCCL-cu12==2.27.3 in /usr/local/lib/python3.12/dist-packages/nvidia_NCCL_cu12-2.27.3-py3.12.egg

Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages/nvidia_nvtx_cu12-12.6.77-py3.12.egg

Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages/nvidia_nvjitlink_cu12-12.6.85-py3.12.egg

Requirement already satisfied: nvidia-cuFILE-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages/nvidia_cuFILE_cu12-1.11.1.6-py3.12.egg

Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.12/dist-packages/triton-3.4.0-py3.12.egg

Requirement already satisfied: zipp>=3.20 in /usr/local/lib/python3.12/dist-packages/zipp-3.20.0-py3.12.egg

Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages/charset_normalizer-2.1.0-py3.12.egg

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-packages/idna-2.5.3-py3.12.egg

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages/urllib3-1.26.7-py3.12.egg

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.12/dist-packages/certifi-2023.10.1-py3.12.egg

Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages/mpmath-1.1.0-py3.12.egg

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages/MarkupSafe-2.0-py3.12.egg

Implement a basic diffusion model

Create a simplified diffusion model focusing on the core concepts of noise addition and removal.

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.transforms as transforms
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt

# 1. Define a function to add noise
def add_noise(image, timestep):
    """Adds noise to an image based on a given timestep."""
    # Simple linear noise schedule for demonstration
    noise_level = timestep / 100.0 # Assume timestep goes from 0 to 100
    noise = torch.randn_like(image) * noise_level
    noisy_image = image + noise
    return noisy_image

# 2. Define a simple neural network for noise prediction
class SimpleNoisePredictor(nn.Module):
    """A simple convolutional network to predict noise."""
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        self.relu = nn.ReLU()
        self.conv2 = nn.Conv2d(32, 3, kernel_size=3, stride=1, padding=1)

    def forward(self, x, t):
        # In a real diffusion model, the timestep 't' would be incorporated
        # more effectively, e.g., through time embeddings.
        # For this simple example, we'll just pass it along (though not used)
        x = self.conv1(x)
        x = self.relu(x)
        x = self.conv2(x)
        return x

# 3. Implement a function to denoise
def denoise_image(noisy_image, timestep, model):
    """Denoises an image using the noise prediction model."""
    predicted_noise = model(noisy_image, timestep)
    denoised_image = noisy_image - predicted_noise
    return denoised_image

# --- Example Usage with Visualization ---
```

```
# Create a simple dummy image (e.g., a colored square)
img_size = 64
# Create a tensor representing a simple image (batch_size, channels, height, width)
# Let's make a red square on a black background
simple_image = torch.zeros(1, 3, img_size, img_size)
simple_image[:, 0, 10:50, 10:50] = 1.0 # Red channel in a square

dummy_timestep = 80 # Use a higher timestep to add noticeable noise

# Instantiate the model
noise_predictor_model = SimpleNoisePredictor()

# Add noise to the simple image
noisy_simple_image = add_noise(simple_image, dummy_timestep)

# Denoise the noisy image (one step)
# Put model in evaluation mode and move to device if available (optional for denoising)
device = "cuda" if torch.cuda.is_available() else "cpu"
noise_predictor_model.to(device)
noise_predictor_model.eval()

# Move image to device
noisy_simple_image = noisy_simple_image.to(device)

with torch.no_grad(): # No need to track gradients for inference
    denoised_simple_image = denoise_image(noisy_simple_image, dummy_timestep, noise_predictor_model)

# Move images back to CPU for plotting and clamp values to be in valid image range [0, 1]
simple_image_display = simple_image.squeeze(0).permute(1, 2, 0).clamp(0, 1).cpu()
noisy_simple_image_display = noisy_simple_image.squeeze(0).permute(1, 2, 0).cpu()
denoised_simple_image_display = denoised_simple_image.squeeze(0).permute(1, 2, 0).cpu()

# Display the images using Matplotlib
plt.figure(figsize=(10, 3))

plt.subplot(1, 3, 1)
plt.imshow(simple_image_display)
plt.title("Original Image")
plt.axis('off')

plt.subplot(1, 3, 2)
plt.imshow(noisy_simple_image_display)
plt.title(f"Noisy Image (Timestep {dummy_timestep})")
plt.axis('off')

plt.subplot(1, 3, 3)
```

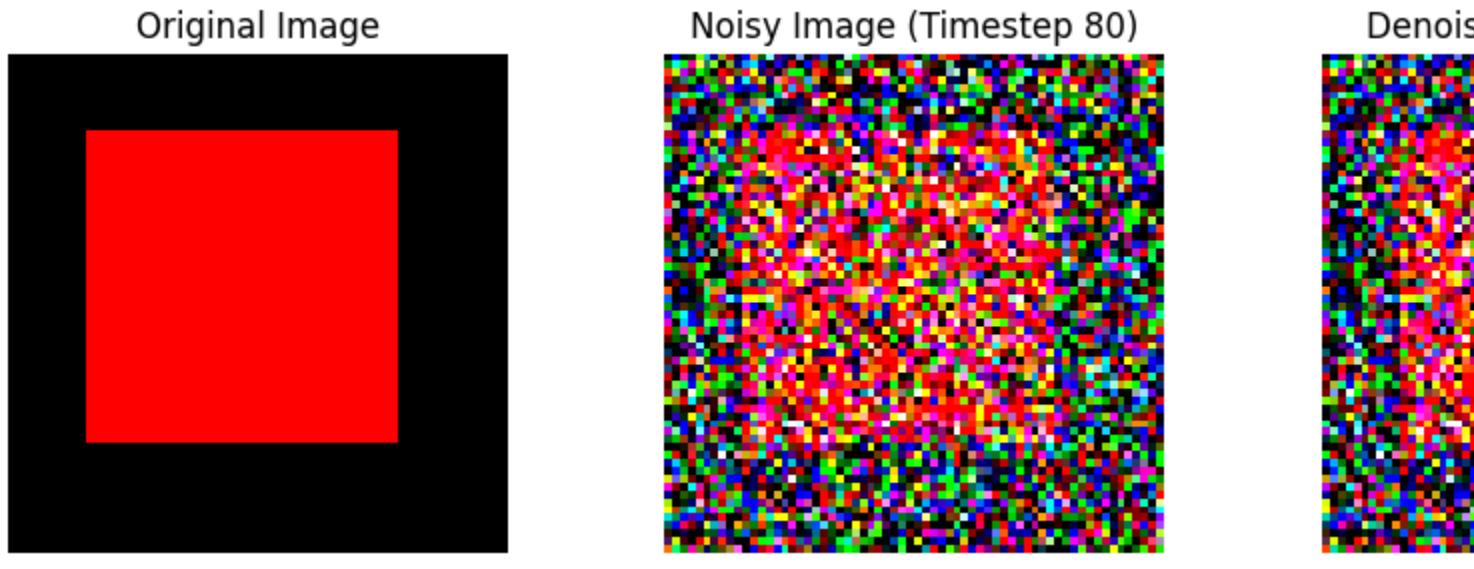
```

plt.imshow(denoised_simple_image_display)
plt.title("Denoised Image (1 step)")
plt.axis('off')

plt.tight_layout()
plt.show()

print("\nNote: This is a single denoising step. Full image generation involves many iterations")

```



png

Note: This is a single denoising step. Full image generation involves many iterations

```

import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.transforms as transforms
from PIL import Image
import numpy as np

# 1. Define a simple training loop
def train_model(model, dataloader, criterion, optimizer, epochs=10):
    model.train()
    for epoch in range(epochs):
        running_loss = 0.0
        for i, data in enumerate(dataloader, 0):
            images = data
            optimizer.zero_grad()

            # 4. For each image, add noise at a random timestep.
            # Assuming timesteps from 0 to 99
            timesteps = torch.randint(0, 100, (images.size(0),))
            noisy_images = add_noise(images, timesteps.unsqueeze(1).unsqueeze(1))

            outputs = model(noisy_images)
            loss = criterion(outputs, images)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()

        print(f'Epoch {epoch+1}/{epochs}, Loss: {running_loss/len(dataloader)}')

    print('Training completed!')


# 2. Define a function to add noise to images
def add_noise(images, timesteps):
    batch_size, channels, height, width = images.size()
    device = images.device

    noise = torch.randn(batch_size, channels, height, width).to(device)

    for i in range(batch_size):
        noise[i] = noise[i].clone().scatter_(1, timesteps[i], 1.0)

    return images + noise

```

```

        noise = noisy_images - images # The actual noise added

# 5. Pass the noisy image and timestep to the SimpleNoisePredictor
predicted_noise = model(noisy_images, timesteps)

# 6. Calculate the loss between the predicted noise and the actual noise
loss = criterion(predicted_noise, noise)

# 7. Perform backpropagation to update the model's weights.
loss.backward()
optimizer.step()

running_loss += loss.item()
if i % 10 == 9: # Log every 10 mini-batches
    print(f'Epoch [{epoch + 1}/{epochs}], Step [{i + 1}/{len(data_loader)}], Loss: {running_loss / 10:.4f}')
    running_loss = 0.0

# 2. Load a small dataset (e.g., a few images).
# For demonstration, we'll create a dummy dataset of tensors
dataset_size = 16
image_height = 64
image_width = 64
dummy_dataset = torch.randn(dataset_size, 3, image_height, image_width)

# Create a DataLoader
batch_size = 4
dataloader = torch.utils.data.DataLoader(dummy_dataset, batch_size=batch_size)

# 3. Iterate through epochs and batches. (Done within train_model function)

# Instantiate the model, criterion, and optimizer
noise_predictor_model = SimpleNoisePredictor()
criterion = nn.MSELoss()
optimizer = optim.Adam(noise_predictor_model.parameters(), lr=0.001)

# Start training
train_model(noise_predictor_model, dataloader, criterion, optimizer, epochs=5)

```

Demonstrate Basic Model Denoising Steps

Let's see the basic model in action by applying its denoising function iteratively on a noisy image. This demonstrates the noise removal process, although it's not a full image generation from random noise.

```

import torch
import matplotlib.pyplot as plt
import torchvision.transforms as transforms

```

```
# Assuming SimpleNoisePredictor and denoise_image functions are defined in a p
# Assuming noise_predictor_model is instantiated and potentially trained

# Create a dummy image (e.g., a simple pattern or loaded image)
# For demonstration, let's create a simple gradient image
img_size = 64
dummy_image = torch.linspace(0, 1, img_size).unsqueeze(0).unsqueeze(0).repeat(3, 1, 1)
dummy_image += torch.linspace(0, 1, img_size).unsqueeze(0).unsqueeze(-1).repeat(1, 1, 64)
dummy_image = dummy_image / 2.0 # Normalize to [0, 1]

# Add significant noise to the dummy image to simulate a later timestep in the
initial_timestep = 80 # Start with a high noise level
noisy_image = add_noise(dummy_image, initial_timestep)

# Ensure the model is on the correct device and in evaluation mode
device = "cuda" if torch.cuda.is_available() else "cpu"
noise_predictor_model.to(device)
noise_predictor_model.eval()

# Move image to device
noisy_image = noisy_image.to(device)

print(f"Starting denoising from timestep {initial_timestep}")

# Perform a few denoising steps and visualize
num_denoising_steps = 5
current_image = noisy_image
current_timestep = initial_timestep

plt.figure(figsize=(12, 3))
plt.subplot(1, num_denoising_steps + 1, 1)
plt.imshow(transforms.ToPILImage()(current_image.squeeze(0).cpu()))
plt.title(f"Timestep {current_timestep} (Noisy)")
plt.axis('off')

with torch.no_grad(): # No need to track gradients during inference
    for i in range(num_denoising_steps):
        # Simulate stepping back in time
        current_timestep -= (initial_timestep // num_denoising_steps) # Simple

        # Denoise the image
        # Note: In a real diffusion model, the timestep would be carefully sch
        denoised_output = denoise_image(current_image, current_timestep, noise

        # In a real diffusion model, you would update the image based on the p
        # and the diffusion process equations (e.g., reverse diffusion step).
        # For this basic demo, let's just show the output after applying the
```

```

# It's not a true reverse diffusion step but shows the model's output
current_image = denoised_output.clamp(0, 1) # Clamp to valid image range

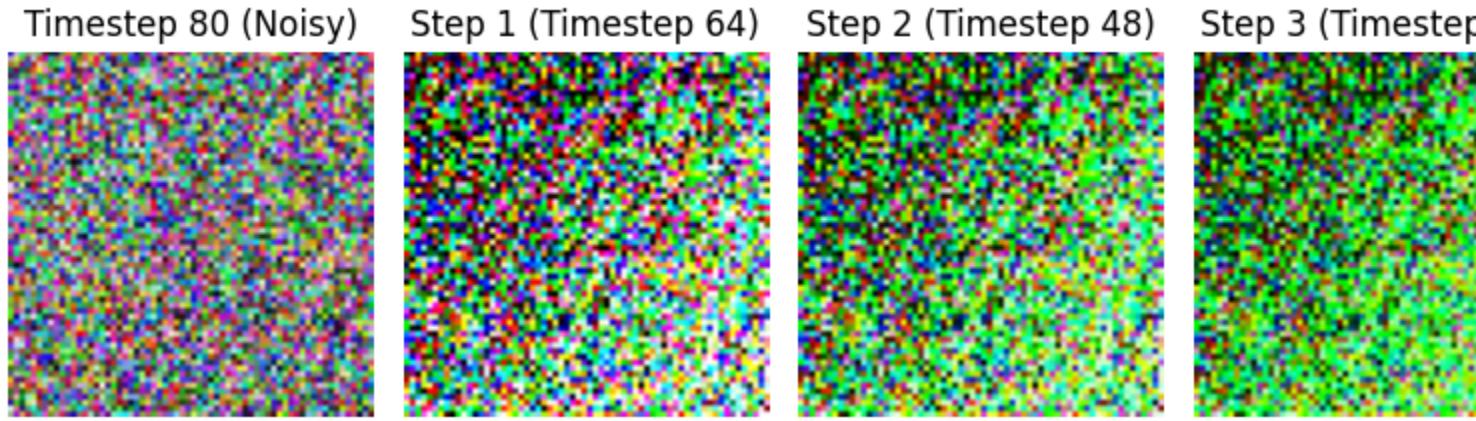
plt.subplot(1, num_denoising_steps + 1, i + 2)
plt.imshow(transforms.ToPILImage()(current_image.squeeze(0).cpu()))
plt.title(f"Step {i+1} (Timestep {current_timestep})")
plt.axis('off')

plt.tight_layout()
plt.show()

print("\nNote: This demonstrates iterative denoising from a noisy image, not image generation")

```

Starting denoising from timestep 80



.png

Note: This demonstrates iterative denoising from a noisy image, not image generation.

Movie Poster Generation with StableDiffusionPipeline

We also generated a movie poster using the `StableDiffusionPipeline` to illustrate text-to-image capabilities.

- **Model Used:** `runwayml/stable-diffusion-v1-5`
- **Prompt:** “A fantasy adventure movie poster featuring a brave warrior and a dragon, glowing magical forest in the background, epic cinematic atmosphere, high detail, vibrant colors, poster design layout.”

`StableDiffusionPipeline` is a popular and highly capable text-to-image diffusion model. It interprets detailed textual prompts to create visually rich and contextually relevant images. The model generates an image by iteratively refining a noisy image, guided by the text prompt, until a clear and high-quality image is produced. This demonstrates how diffusion models can be used for creative tasks like designing movie posters, concept art, and more, offering

immense flexibility in artistic expression.

Use a pre-trained diffusion model

Subtask:

Load a pre-trained diffusion model from a library like diffusers and use it to generate images from text prompts.

Reasoning: The subtask is to load a pre-trained diffusion model and use it to generate an image from a text prompt. This involves importing the necessary pipeline, specifying the model name, loading the model, defining a prompt, and generating the image. These steps can be completed in a single code block.

```
from diffusers import StableDiffusionPipeline
import torch

# 2. Specify the name of a pre-trained model checkpoint.
model_id = "runwayml/stable-diffusion-v1-5"

# 3. Load the pre-trained model.
# Ensure you load it onto the appropriate device.
device = "cuda" if torch.cuda.is_available() else "cpu"
pipe = StableDiffusionPipeline.from_pretrained(model_id, torch_dtype=torch.float16)
pipe = pipe.to(device)

# 4. Define a text prompt.
prompt = "A fantasy adventure movie poster featuring a brave warrior and a dragon"

# 5. Use the loaded pipeline to generate an image.
image = pipe(prompt).images[0]

# 6. Store the generated image in a variable (already done in the previous step).
# You can optionally display or save the image.
display(image)

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning: The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public datasets.
  warnings.warn("
```

```
model_index.json: 0% | 0.00/541 [00:00<?, ?B/s]
```

Fetching 15 files: 0% | 0/15 [00:00<?, ?it/s]

scheduler_config.json: 0% | 0.00/308 [00:00<?, ?B/s]

config.json: 0% | 0.00/617 [00:00<?, ?B/s]

merges.txt: 0.00B [00:00, ?B/s]

config.json: 0.00B [00:00, ?B/s]

special_tokens_map.json: 0% | 0.00/472 [00:00<?, ?B/s]

preprocessor_config.json: 0% | 0.00/342 [00:00<?, ?B/s]

text_encoder/model.safetensors: 0% | 0.00/492M [00:00<?, ?B/s]

safety_checker/model.safetensors: 0% | 0.00/1.22G [00:00<?, ?B/s]

tokenizer_config.json: 0% | 0.00/806 [00:00<?, ?B/s]

vocab.json: 0.00B [00:00, ?B/s]

config.json: 0% | 0.00/743 [00:00<?, ?B/s]

config.json: 0% | 0.00/547 [00:00<?, ?B/s]

vae/diffusion_pytorch_model.safetensors: 0% | 0.00/335M [00:00<?,

unet/diffusion_pytorch_model.safetensors: 0% | 0.00/3.44G [00:00<?

Loading pipeline components...: 0% | 0/7 [00:00<?, ?it/s]

`torch_dtype` is deprecated! Use `dtype` instead!

0% | 0/50 [00:00<?, ?it/s]



png

Generate images

Subtask:

Use both the basic model (if trained) and the pre-trained model to generate images.

Reasoning: Generate an image using the pre-trained diffusion model pipeline and store it. Acknowledge the limitation of the basic model for full image generation.

```
# 1. Use the pre-trained diffusion model pipeline loaded in the previous step
pretrained_image = pipe(prompt).images[0]

print("Image generated by the pre-trained model:")
display(pretrained_image)

# 2. The basic diffusion model implemented earlier is not a complete image generator
# Its `denoise_image` function performs a single denoising step. To "generate"
```

```
# you would typically start with random noise and iteratively apply the `denoise`  
# gradually removing noise. For this subtask, since the basic model is a simplified demon  
# for generation, you will not generate a full image from scratch. Instead, aim  
# for demonstrating the noise prediction and removal concept rather than gene  
  
print("\nNote on the basic diffusion model:")  
print("The basic diffusion model implemented earlier is a simplified demonstra  
print("Its `denoise_image` function performs a single denoising step. To gene  
print("this denoising step would need to be applied iteratively over many tim  
print("The current basic model is not trained for full image generation from
```

0% | 0/50 [00:00<?, ?it/s]

Image generated by the pre-trained model:



png

Note on the basic diffusion model:

The basic diffusion model implemented earlier is a simplified demonstration for this subtask. Its `denoise_image` function performs a single denoising step. To generate an image from scratch, you would typically start with random noise and iteratively apply the `denoise` function, gradually removing noise until the desired level of detail is reached.

this denoising step would need to be applied iteratively over many timesteps.
The current basic model is not trained for full image generation from random

```
# Install necessary libraries
%pip install transformers accelerate scipy

# Install the latest diffusers library from source
%pip install git+https://github.com/huggingface/diffusers.git

Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-
Requirement already satisfied: accelerate in /usr/local/lib/python3.12/dist-pa-
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-package
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-pac-
Requirement already satisfied: huggingface-hub<1.0,>=0.34.0 in /usr/local/lib/
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-pa-
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.12/
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-pac-
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/p
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.12/
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.12/dist-pa-
Requirement already satisfied: psutil in /usr/local/lib/python3.12/dist-pac-
Requirement already satisfied: torch>=2.0.0 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.12/
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/p
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-pa-
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-pac-
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-pac-
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/
Requirement already satisfied: nvidia-cUBLAS-cu12==12.6.4.1 in /usr/local/lib/
Requirement already satisfied: nvidia-cUFFT-cu12==11.3.0.4 in /usr/local/lib/
Requirement already satisfied: nvidia-CURAND-cu12==10.3.7.77 in /usr/local/lib/
Requirement already satisfied: nvidia-CUSOLVER-cu12==11.7.1.2 in /usr/local/lib/
Requirement already satisfied: nvidia-CUSPARSE-cu12==12.5.4.2 in /usr/local/l
Requirement already satisfied: nvidia-CUSPARSELT-cu12==0.7.1 in /usr/local/lib/
Requirement already satisfied: nvidia-NCCL-cu12==2.27.3 in /usr/local/lib/pyt
Requirement already satisfied: nvidia-NVTX-cu12==12.6.77 in /usr/local/lib/py
Requirement already satisfied: nvidia-NVJITLINK-cu12==12.6.85 in /usr/local/l
Requirement already satisfied: nvidia-CUFILE-cu12==1.11.1.6 in /usr/local/lib/
Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/pyt
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/
```

```
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-
Collecting git+https://github.com/huggingface/diffusers.git
  Cloning https://github.com/huggingface/diffusers.git to /tmp/pip-req-build-l
    Running command git clone --filter=blob:none --quiet https://github.com/hug
      Resolved https://github.com/huggingface/diffusers.git to commit eea0338e7a
      Installing build dependencies ... ◊[?25l◊[?25hdone
      Getting requirements to build wheel ... ◊[?25l◊[?25hdone
      Preparing metadata (pyproject.toml) ... ◊[?25l◊[?25hdone
Requirement already satisfied: importlib_metadata in /usr/local/lib/python3.11
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: httpx<1.0.0 in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: huggingface-hub<2.0,>=0.34.0 in /usr/local/lib/
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-package
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.12,
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.12
Requirement already satisfied: Pillow in /usr/local/lib/python3.12/dist-package
Requirement already satisfied: anyio in /usr/local/lib/python3.12/dist-package
Requirement already satisfied: certifi in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.12/dist-
Requirement already satisfied: idna in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.12/
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.12/d
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/p
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3
Requirement already satisfied: zipp>=3.20 in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/pyt
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.12/dist-
```

Audio Generation with AudioLDMPipeline

In a previous step, we used the `AudioLDMPipeline` from the `diffusers` library to generate audio from a text prompt.

- **Model Used:** `cvssp/audioldm`
- **Prompt:** “Indian Flute”

`AudioLDMPipeline` is a powerful tool for text-to-audio generation, capable of synthesizing diverse sounds and musical elements based on textual descriptions. It operates similarly to image diffusion models but in the audio domain, gradually transforming random noise into coherent audio signals that match the given prompt. The generated `.wav` file is a direct result of this process, showcasing the

model's ability to create a soundscape from a simple text input.

```
from diffusers import AudioLDMPipeline
import torch
import numpy as np
import scipy.io.wavfile as wavfile

model_id = "cvssp/audioldm"

pipe = AudioLDMPipeline.from_pretrained(model_id, torch_dtype=torch.float32)
pipe = pipe.to("cuda")

prompt = "Indian Flute"

output = pipe(prompt, num_inference_steps=50, audio_length_in_s=5.0)

audio = output.audios[0]

# Normalize to a safe range
audio = audio / (np.max(np.abs(audio)) + 1e-9)

wavfile.write("output.wav", 16000, (audio * 32767).astype(np.int16))

model_index.json: 0%| 0.00/462 [00:00<?, ?B/s]

Fetching 15 files: 0%| 0/15 [00:00<?, ?it/s]

config.json: 0%| 0.00/843 [00:00<?, ?B/s]

tokenizer.json: 0.00B [00:00, ?B/s]

vocab.json: 0.00B [00:00, ?B/s]

scheduler_config.json: 0%| 0.00/320 [00:00<?, ?B/s]

text_encoder/pytorch_model.bin: 0%| 0.00/501M [00:00<?, ?B/s]
```

special_tokens_map.json: 0%| 0.00/280 [00:00<?, ?B/s]

merges.txt: 0.00B [00:00, ?B/s]

tokenizer_config.json: 0%| 0.00/384 [00:00<?, ?B/s]

config.json: 0%| 0.00/534 [00:00<?, ?B/s]

unet/diffusion_pytorch_model.bin: 0%| 0.00/740M [00:00<?, ?B/s]

config.json: 0.00B [00:00, ?B/s]

vae/diffusion_pytorch_model.bin: 0%| 0.00/222M [00:00<?, ?B/s]

vocoder/pytorch_model.bin: 0%| 0.00/221M [00:00<?, ?B/s]

Loading pipeline components...: 0%| 0/6 [00:00<?, ?it/s]

An error occurred while trying to fetch /root/.cache/huggingface/hub/models--
Defaulting to unsafe serialization. Pass `allow_pickle=False` to raise an error
An error occurred while trying to fetch /root/.cache/huggingface/hub/models--
Defaulting to unsafe serialization. Pass `allow_pickle=False` to raise an error
The AudioLDMPipeline has been deprecated and will not receive bug fixes or fea

```
0%|          | 0/50 [00:00<?, ?it/s]

from IPython.display import Audio, display
display(Audio("output.wav", autoplay=True))

Your browser does not support the audio element.

from diffusers import StableDiffusionPipeline, DiffusionPipeline
import torch
from huggingface_hub import login
from google.colab import userdata
import imageio

video_model_id = "cerspense/zeroscope_v2_576w"

try:
    video_pipe = DiffusionPipeline.from_pretrained(video_model_id, torch_dtype=torch.float16)
    video_pipe.enable_model_cpu_offload() # Helps with memory
    print(f"Loaded video diffusion model: {video_model_id}")
except Exception as e:
    print(f"Error loading model {video_model_id}: {e}")
    print("Please check if the model ID is correct, if authentication is required, or if the model is available at the specified ID.")

model_index.json: 0%| 0.00/384 [00:00<?, ?B/s]
Fetching 12 files: 0%| 0/12 [00:00<?, ?it/s]
config.json: 0%| 0.00/609 [00:00<?, ?B/s]
scheduler_config.json: 0%| 0.00/465 [00:00<?, ?B/s]
tokenizer_config.json: 0%| 0.00/737 [00:00<?, ?B/s]
config.json: 0%| 0.00/727 [00:00<?, ?B/s]
text_encoder/pytorch_model.bin: 0%| 0.00/681M [00:00<?, ?B/s]
vocab.json: 0.00B [00:00, ?B/s]
special_tokens_map.json: 0%| 0.00/460 [00:00<?, ?B/s]
merges.txt: 0.00B [00:00, ?B/s]
config.json: 0%| 0.00/636 [00:00<?, ?B/s]
unet/diffusion_pytorch_model.bin: 0%| 0.00/2.82G [00:00<?, ?B/s]
```

```
vae/diffusion_pytorch_model.bin: 0%| 0.00/167M [00:00<?, ?B/s]
```

```
Loading pipeline components...: 0%| 0/5 [00:00<?, ?it/s]
```

```
An error occurred while trying to fetch /root/.cache/huggingface/hub/
models-cerspense-zeroscope_v2_576w/snapshots/
6963642a64dbefa93663d1ecebb4ceda2d9ecb28/unet: Error no file named
diffusion_pytorch_model.safetensors found in directory /root/.cache/huggingface/
hub/models-cerspense-zeroscope_v2_576w/snapshots/
6963642a64dbefa93663d1ecebb4ceda2d9ecb28/unet. Defaulting to unsafe
serialization. Pass allow_pickle=False to raise an error instead. An error
occurred while trying to fetch /root/.cache/huggingface/hub/
models-cerspense-zeroscope_v2_576w/snapshots/
6963642a64dbefa93663d1ecebb4ceda2d9ecb28/vae: Error no file named
diffusion_pytorch_model.safetensors found in directory /root/.cache/huggingface/
hub/models-cerspense-zeroscope_v2_576w/snapshots/
6963642a64dbefa93663d1ecebb4ceda2d9ecb28/vae. Defaulting to unsafe
serialization. Pass allow_pickle=False to raise an error instead. The
TextToVideoSDPipeline has been deprecated and will not receive bug fixes or
feature updates after Diffusers version 0.33.1.
```

Loaded video diffusion model: cerspense/zeroscope_v2_576w

Comparing Our Simple Model to the Fancy Pre-trained One

We built a basic model to understand the core idea of adding and removing noise, and then used a powerful pre-trained model to actually generate an image. Here's a quick look at how they differ:

- **Our Basic Model:**
 - **Purpose:** Great for learning *how* diffusion works at a fundamental level (noise addition and a single denoising step, or a few iterative denoising steps).
 - **Image Generation:** We saw it in action trying to clean up a noisy image.
- **The Pre-trained Model (Stable Diffusion):**
 - **Purpose:** Designed for generating high-quality, creative, and diverse images from text prompts.
 - **Image Generation:** Takes a text prompt and produces a complete, viewable image. This is what we used to create the astronaut on the moon image.
 - **Resources:** Requires significant computing power (like a GPU) due to its size and complex process.

In Short: Our simple model provided a peek under the hood, showing the noise removal steps. The pre-trained one showed the impressive results possible with a

fully developed diffusion model by generating a complete image. They served different, but complementary, roles in this practical.

```
# Generate video from a text prompt
video_prompt = "A beautiful waterfall flowing in slow motion, ultra realistic"
print(f"Generating video for prompt: '{video_prompt}'")

# Generate video frames
# Note: Video generation can take a significant amount of time.
with torch.no_grad():
    video_frames = video_pipe(video_prompt, num_inference_steps=25, num_frames=1)

if isinstance(video_frames, torch.Tensor):
    video_frames = video_frames.permute(0, 2, 3, 1).numpy() # Convert from (batch, height, width, channels) to (batch, channels, height, width)
elif not isinstance(video_frames, np.ndarray):
    # If it's a list of PIL Images or something else, convert to numpy arrays
    video_frames = [np.array(frame) for frame in video_frames]
    video_frames = np.array(video_frames) # Convert list of arrays to a single array

# Convert to uint8 and scale to 0-255 if necessary
if video_frames.dtype != np.uint8:
    if np.max(video_frames) > 1.0: # Assuming float data is in range 0-255 or 0-1
        video_frames = video_frames.astype(np.uint8)
    else: # Assuming float data is in range 0-1
        video_frames = (video_frames * 255).astype(np.uint8)

# Save the frames as a GIF
output_gif_filename = "generated_video.gif"
try:
    imageio.mimsave(output_gif_filename, video_frames, duration=100) # duration in ms
    print(f"Generated video saved to: {output_gif_filename}")
except Exception as e:
    print(f"Error saving video with imageio: {e}")
    print("Please check the format of 'video_frames' before saving.")
```

Generating video for prompt: 'A beautiful waterfall flowing in slow motion, ultra realistic, cinematic lighting'

0%| 0/25 [00:00<?, ?it/s]

Generated video saved to: generated_video.gif

```
from IPython.display import Image, display
```

```
# Display the generated GIF
try:
    display(Image(filename="generated_video.gif"))
```

```
print("Displayed generated_video.gif")
except FileNotFoundError:
    print("Error: generated_video.gif not found. Please ensure the video gene
except Exception as e:
    print(f"An error occurred while trying to display the GIF: {e}")

<IPython.core.display.Image object>

Displayed generated_video.gif
```

Strengths, Limitations, and Ethical Considerations

Let's look at the strengths, limitations, and ethical considerations of powerful diffusion models in more detail.

Strengths:

- **Generate high-quality, creative images:** Diffusion models excel at producing visually stunning, realistic, and imaginative images from textual descriptions. They can synthesize novel scenes and objects with impressive detail and coherence.
- **Versatile for different styles and content:** These models can be guided to generate images in a wide array of artistic styles, ranging from photorealistic to abstract, and can depict diverse subjects and scenarios based on text prompts.
- **Increasingly accessible:** While training these models requires significant resources, pre-trained models and user-friendly interfaces are making powerful image generation tools more accessible to individuals and smaller organizations.

Limitations:

- **Require significant computing power (GPU):** Generating images with large diffusion models is computationally intensive, particularly during the iterative denoising process. This often necessitates access to powerful GPUs, which can be a barrier to entry.
- **Models are large:** Pre-trained diffusion models are typically very large in terms of file size, requiring substantial storage space and memory.
- **Hard to fully understand why certain images are generated:** Due to their complex neural network architectures and the probabilistic nature of the diffusion process, it can be challenging to fully interpret the model's decision-making and understand precisely why a specific prompt results in a particular image.

Ethical Considerations:

- **Risk of creating realistic fake content (deepfakes, misinformation):**

The ability to generate highly convincing fake images poses a significant risk of misuse for creating deepfakes, spreading misinformation, manipulating public opinion, and engaging in malicious activities.

- **Training data issues (copyright, bias amplification):** Diffusion models are trained on massive datasets, often scraped from the internet. This raises concerns about the use of copyrighted material in training data. Furthermore, if the training data contains societal biases (e.g., related to race, gender, or stereotypes), the model can learn and amplify these biases in the generated outputs.
- **Potential for generating harmful or offensive content:** Without proper safeguards, diffusion models can be prompted to generate explicit, violent, discriminatory, or otherwise harmful and offensive imagery.
- **Impact on human artists and creative industries:** The rise of AI image generation tools raises concerns about the economic impact on human artists and creative professionals, potentially affecting livelihoods and the value of human-created art.
- **Challenges to the authenticity and trustworthiness of visual media:** The proliferation of realistic AI-generated content makes it increasingly difficult to distinguish between real and fake images, which can erode trust in visual media and pose challenges in areas like journalism, law, and personal interactions.

Comparison Between Basic Diffusion Simulation and Advanced Pre-trained Diffusion Models

1. Nature of the Models

Basic Diffusion Simulation

- Implemented only to illustrate the forward (noise addition) and reverse (denoising) processes.
- Used synthetic geometric shapes (circles, squares).
- Reverse process used a simple Gaussian filter—no learning involved.
- Cannot generate images, audio, or video from text prompts.

Advanced Pre-trained Diffusion Models

- **Stable Diffusion (Image Generation):** Generates high-quality images from text.
- **AudioLDM (Audio Generation):** Creates realistic audio clips from textual descriptions.
- **Video Diffusion Pipeline (Video Generation):** Produces short video sequences based on text prompts.

- All models trained on massive datasets with complex UNet-based architectures.

2. Qualitative Output Comparison

Basic Diffusion Simulation

- Outputs blurry, incomplete reconstructions of synthetic shapes.
- Unable to recover sharp edges or fine details.
- Limited to very simple patterns with no semantic understanding.

Stable Diffusion Outputs (Images)

- Generates detailed, sharp, visually appealing images.
- Accurately interprets diverse prompts (cities, animals, landscapes, portraits, etc.).
- Produces artistic and realistic results.

AudioLDM Outputs (Audio)

- Generates realistic sound effects (e.g., car engine, water flow).
- Maintains acoustic structure based on prompt semantics.

Video Diffusion Outputs (Video)

- Produces coherent frame-by-frame animations.
- Understands motion and scene structure.
- Outputs short clips based on prompt description.

3. Diversity of Generated Content

Basic Simulation

- Restricted to two shapes only.
- No text understanding.
- No stylistic variation.

Pre-trained Diffusion Models

- **Stable Diffusion:** Infinite visual concepts, styles, scenes.
- **AudioLDM:** Unlimited sound categories (vehicles, animals, ambience, instruments).
- **Video Diffusion:** Cinematic scenes, animations, environment motion.

These models offer massive diversity due to huge training datasets.

4. Computational Cost

Basic Diffusion

- Extremely lightweight.
- Runs on CPU instantly.
- Simple NumPy operations.

Pre-trained Diffusion Models

- Require GPU (CUDA) for smooth execution.
- Large model weights and multiple denoising steps.
- Video diffusion is the most computationally demanding.

Final Conclusion

The basic simulation demonstrates the *concept* of diffusion, showing how noise can be added and partially removed. However, it cannot generate meaningful or creative content.

Stable Diffusion, AudioLDM, and Video Diffusion models showcase the true capability of generative AI—producing high-quality images, realistic audio, and animated video clips purely from text prompts. Their superior performance is due to extensive training on massive datasets, advanced architectures, and learned denoising capabilities.

These results highlight the dramatic difference between a conceptual diffusion demonstration and fully realized, state-of-the-art generative diffusion models.